

Roman-empire: MLP, depth 9, hidden width 512, dropout 0.425, norm batch, residual True
Alignment + REPAIR comparison, 3 endpoint pair(s)

